

MPI

Spring Training 2025

Yoshihiro Izawa

2025/05/13

Contents

| | |
|----------------------------------------------------|----|
| 1. MPI Overview | 3 |
| 1.1 What is MPI (Message Passing Interface)? | 4 |
| 1.2 Parallel Programming Classification | 5 |
| 1.3 MPI Features | 6 |
| 1.4 Typical example of Usage | 7 |
| 1.5 Comparison between implementations | 8 |
| 1.6 Key Communication Primitives | 9 |
| 2. Basic Learning of MPI | 10 |
| 2.1 Minimum MPI Program | 11 |
| 2.2 Important Terms of MPI | 15 |
| 2.3 Point-to-Point Communication | 21 |
| 3. Collective Communication | 28 |
| 3.1 synchronization | 29 |
| 3.2 MPI_Barrier | 30 |
| 3.3 MPI_Bcast | 32 |

| | |
|--------------------------------------------------|----|
| 3.4 MPI_Scatter, MPI_Gather, MPI_Allgather | 37 |
| 3.5 MPI_Reduce, MPI_Allreduce | 43 |
| 4. References | 48 |
| 4.1 MPI Reference | 49 |

1. MPI Overview

1.1 What is MPI (Message Passing Interface)?

- **A standard API** for message passing between distributed memories in parallel computing.
- MPI assumes a **distributed-memory computing system**
- MPI can run on **shared-memory computing system**
- MPI programming model (basically) uses **SIMD**

1.2 Parallel Programming Classification

- **Multi-Process**: MPI(Message Passing Interface), HPF(High Performance Fortran)
- **Multi-Thread**: OpenMP, Pthread(POSIX Thread)

1.3 MPI Features

- **Communication Model:**
 - Uses message passing for communication between processes.
- **Distributed Memory Support:**
 - Each process has its own memory space, no shared memory.
- **Multi-node Capacity:**
 - Can run across multiple nodes; abstracts network communication.
- **Standardized API:**
 - Standardized interface in C, C++, and Fortran; highly portable.
- **Multiple Implementation:**
 - Available implementations include OpenMPI, MPICH, and Intel MPI, etc.
- **Difficult to Debug:**
 - Debugging is challenging due to concurrency and communication complexity.

1.4 Typical example of Usage

- Simulation on a supercomputer(Physics, Meteorology, Chemistry, etc.)
- Data processing in large-scale data analysis (e.g., genomics, astronomy).
- Machine learning training on large datasets (e.g., distributed deep learning).

1.5 Comparison between implementations

| | OpenMPI | MPICH | Intel MPI |
|---------------------|-----------------------------------|----------------------------------|-------------------------------------|
| Developer | Universities, Companies | Argonne National Laboratory | Intel Corporation |
| Distribution | Open source | Open source | Free version included |
| Optimization Target | General purpose | Lightweight, stable | Optimized for Intel architecture |
| Performance | Medium to high | Lightweight, stable, scalable | Best performance on Intel CPUs |
| Main Use | Academic clusters, general HPC | Research, education | Commercial HPC, Intel clusters |

1.6 Key Communication Primitives

- **System function:** MPI_Init, MPI_Finalize, MPI_Comm_size, MPI_Comm_rank
- **Point-to-point communication:** MPI_Send, MPI_Recv
- **Collective communication:** MPI_Bcast, MPI_Reduce, MPI_Alltoall
- **Synchronization:** MPI_Barrier, MPI_Wait, MPI_Test
- **Derived data types:** MPI_Type_create_struct, MPI_Type_vector
- **Non-blocking communication:** MPI_Isend, MPI_Irecv
- **Remote memory access:** MPI_Put, MPI_Get
- **Process management:** MPI_Comm_spawn, MPI_Comm_free

2. Basic Learning of MPI

2.1 Minimum MPI Program

2.1.1 Hello World (C)

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    int num_procs;
    int my_rank;

    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    printf("Num of Proc : %d\n", num_procs);
    printf("My Rank      : %d\n", my_rank);

    MPI_Finalize();
    return EXIT_SUCCESS;
}
```

```
mpicc mpi_hello.c -o mpi_hello
mpirun -np 4 ./mpi_hello
Num of Proc : 4
My Rank      : 3
Num of Proc : 4
My Rank      : 2
Num of Proc : 4
My Rank      : 0
Num of Proc : 4
My Rank      : 1
```

2.1 Minimum MPI Program

2.1.2 Hello World (C++)

```
#include <mpi.h>
#include <iostream>
#include <cstdlib>

int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    int num_procs;
    int my_rank;

    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    std::cout << "Num of Proc : " << num_procs <<
std::endl;
    std::cout << "My Rank      : " << my_rank << std::endl;

    MPI_Finalize();
    return EXIT_SUCCESS;
}
```

```
mpic++ mpi_hello.cpp -o mpi_hello
mpirun -np 4 ./mpi_hello
Num of Proc : 4
My Rank      : 3
Num of Proc : 4
My Rank      : 1
Num of Proc : 4
My Rank      : 0
Num of Proc : 4
My Rank      : 2
```

2.1 Minimum MPI Program

2.1.3 Hello World (Fortran)

```
program hello_mpi
  use mpi
  implicit none

  integer :: ierr, rank, size

  call MPI_Init(ierr)
  call MPI_Comm_rank(MPI_COMM_WORLD, rank, ierr)
  call MPI_Comm_size(MPI_COMM_WORLD, size, ierr)

  print *, "Num of Proc:", size
  print *, "My Rank:    ", rank

  call MPI_Finalize(ierr)
end program hello_mpi
```

```
mpif90 mpi_hello.f90 -o mpi_hello
mpirun -np 4 ./mpi_hello
Num of Proc: 4
My Rank:    2
Num of Proc: 4
My Rank:    0
Num of Proc: 4
My Rank:    3
Num of Proc: 4
My Rank:    1
```

2.1 Minimum MPI Program

2.1.4 MPI Language Differences

| | C | C++ (※) | Fortran |
|----------------------|-------------------------------------|--------------------------------------------|-------------------------------------------------------|
| MPI Header | <code>#include <mpi.h></code> | <code>#include <mpi.h></code> | <code>use mpi</code> or <code>include 'mpif.h'</code> |
| Official MPI support | ○ | ▲ | ○ |
| Syntax intuitiveness | Explicit C syntax | Almost same as C | <code>call</code> and subroutine based |
| Compiler | <code>mpicc</code> | <code>mpicxx</code> or <code>mpic++</code> | <code>mpif90</code> or <code>mpifort</code> |
| Scientific computing | ○ | ▲ | ◎ |

- C++
 - MPI-3.0 abolished C++ only bindings.
 - Currently, C++ also uses C interface.
- Fortran
 - Considering readability, type safety, and portability, `use mpi` is recommended.

2.2 Important Terms of MPI

2.2.1 Overview

- **Process:**
 - computing unit in parallel computing in MPI.
 - process num is determined by `mpirun -np`
- **Group:**
 - a set of processes that can communicate with each other.
- **Communicator:**
 - a group of processes that can communicate with each other.
- **Rank:**
 - unique identifier for each process in MPI.
 - Ranks are assigned from 0 to `num_procs - 1`.

2.2 Important Terms of MPI

2.2.2 Communicator Image

- Each process belongs to some **group**.
- A group is associated with a **communicator**.
- Each process in a communicator has a unique **rank**.

```
- Example:  
  - Process 0, 1, 2, 3 belong to a group.  
  - Communicator: MPI_COMM_WORLD  
  - Group: [P0, P1, P2, P3]  
  - Rank: 0, 1, 2, 3
```

```
Communicator: MPI_COMM_WORLD  
Group: [P0, P1, P2, P3]  
Rank:  0    1    2    3
```

2.2 Important Terms of MPI

2.2.3 MPI Functions for Communicator

- **MPI_COMM_WORLD:**
 - the default communicator that includes all processes.
 - all processes first belong to this communicator.
 - becomes the default communicator for most MPI functions.
- **MPI_Comm_rank:**
 - retrieves the rank of the calling process in specified communicator.
 - usually use MPI_COMM_WORLD as the communicator.
- **MPI_Comm_size:**
 - retrieves the number of processes in the specified communicator.
 - usually use MPI_COMM_WORLD as the communicator.
- **MPI_Comm_split:**
 - creates a new communicator by splitting the existing one based on a color and key.
 - in other words, create a new communicator with the same color processes.

2.2 Important Terms of MPI

- **MPI_Comm_rank** and **MPI_Comm_size**

```
int rank;  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
printf("I am process %d\n", rank);
```

- **MPI_Comm_size**

```
int size;  
MPI_Comm_size(MPI_COMM_WORLD, &size);  
printf("There are %d processes\n", size);
```

2.2 Important Terms of MPI

- **MPI_Comm_split**

```
// ランクの3の剰余を基に color=0, 1, 2 に分ける
int color = rank % 3;

MPI_Comm new_comm;
MPI_Comm_split(MPI_COMM_WORLD, color, rank, &new_comm);

int new_rank, new_size;
MPI_Comm_rank(new_comm, &new_rank);
MPI_Comm_size(new_comm, &new_size);

printf("World Rank %d => Group %d, New Rank %d of %d\n",
       rank, color, new_rank, new_size);

MPI_Comm_free(&new_comm); // 新しいコミュニケータの解放
```

```
$ mpicc comm_split.c -o comm_split
$ mpirun -np 8 ./comm_split
```

```
World Rank 5 => Group 2, New Rank 1 of 2
World Rank 2 => Group 2, New Rank 0 of 2
World Rank 1 => Group 1, New Rank 0 of 3
World Rank 4 => Group 1, New Rank 1 of 3
World Rank 7 => Group 1, New Rank 2 of 3
World Rank 3 => Group 0, New Rank 1 of 3
World Rank 6 => Group 0, New Rank 2 of 3
World Rank 0 => Group 0, New Rank 0 of 3
```

2.2 Important Terms of MPI

| World Rank | color(= rank % 3) | new group | new rank | new_size |
|------------|-------------------|-----------|----------|----------|
| 0 | 0 | {0, 3, 6} | 0 | 3 |
| 1 | 1 | {1, 4, 7} | 0 | 3 |
| 2 | 2 | {2, 5} | 0 | 2 |
| 3 | 0 | {0, 3, 6} | 1 | 3 |
| 4 | 1 | {1, 4, 7} | 1 | 3 |
| 5 | 2 | {2, 5} | 1 | 2 |
| 6 | 0 | {0, 3, 6} | 2 | 3 |
| 7 | 1 | {1, 4, 7} | 2 | 3 |

2.3 Point-to-Point Communication

send_data()

```
int send_data[10];
for (int i = 0; i < 10; i++)
    send_data[i] = i + 1;
int data_count = 10;

printf("Rank 0: Sending data.\n");
printf("send_data: [");
for (int i = 0; i < 10; i++)
    printf(" %d", send_data[i]);
printf(" ]\n");

MPI_Send((void*)send_data, data_count, MPI_INT,
1, 0, MPI_COMM_WORLD);
```

recv_data()

```
int data[10];
int data_count = 10;
MPI_Status st;

printf("Rank 1: Receiving data.\n");
MPI_Recv((void*)data, data_count, MPI_INT, 0, 0,
MPI_COMM_WORLD, &st);
printf("recv_data: [");
for (int i = 0; i < 10; i++)
    printf(" %d", data[i]);
printf(" ]\n");
```

2.3 Point-to-Point Communication

```
$ mpicc send_recv.c -o mpi_send_recv  
$ mpirun -np 2 ./mpi_send_recv
```

```
Rank 0: Sending data.
```

```
send_data: [ 1 2 3 4 5 6 7 8 9 10 ]
```

```
Rank 1: Receiving data.
```

```
recv_data: [ 1 2 3 4 5 6 7 8 9 10 ]
```

2.3 Point-to-Point Communication

- We can combine `send_data()` and `recv_data()` into a single program.
- The program can be run with `mpirun -np 2 ./mpi_send_recv2`

```
int data_count = 10;
int numbers[data_count];
if (world_rank == 0) {
    for (int i = 0; i < 10; i++) numbers[i] = i + 1;
    printf("Send Data:");
    for (int i = 0; i < 10; i++)
        printf(" %d", numbers[i]);
    printf("\n");
    MPI_Send((void *)&numbers, data_count, MPI_INT, 1, 0, MPI_COMM_WORLD);
} else if (world_rank == 1) {
    MPI_Recv(&numbers, data_count, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Received Data:");
    for (int i = 0; i < 10; i++)
        printf(" %d", numbers[i]);
    printf("\n");
}
```


2.3 Point-to-Point Communication

```
$ mpicc send_recv2.c -o mpi_send_recv2  
$ mpirun -np 2 ./mpi_send_recv2
```

```
Send Data: 1 2 3 4 5 6 7 8 9 10
```

```
Received Data: 1 2 3 4 5 6 7 8 9 10
```

2.3 Point-to-Point Communication

- many MPI functions have the following signature:

```
MPI_Send(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int destination,  
    int tag,  
    MPI_Comm communicator  
);
```

```
MPI_Recv(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int source,  
    int tag,  
    MPI_Comm communicator,  
    MPI_Status* status  
);
```

2.3 Point-to-Point Communication

- Many MPI functions have the following signature

```
MPI_Send(  
    void* data,           // data buffer address  
    int count,           // number of elements in the buffer  
    MPI_Datatype datatype, // data type of the elements  
    int destination,     // destination process rank  
    int tag,             // message tag (for filtering)  
    MPI_Comm communicator // communicator  
);
```

```
MPI_Recv(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int source,  
    int tag,  
    MPI_Comm communicator,  
    MPI_Status* status  
);
```

2.3 Point-to-Point Communication

| MPI Data Type | C Type |
|------------------------|------------------------|
| MPI_SHORT | short int |
| MPI_INT | int |
| MPI_LONG | long int |
| MPI_LONG_LONG | long long int |
| MPI_UNSIGNED_CHAR | unsigned char |
| MPI_UNSIGNED_SHORT | unsigned short int |
| MPI_UNSIGNED | unsigned int |
| MPI_UNSIGNED_LONG | unsigned long int |
| MPI_UNSIGNED_LONG_LONG | unsigned long long int |
| MPI_FLOAT | float |
| MPI_DOUBLE | double |
| MPI_LONG_DOUBLE | long double |
| MPI_BYTE | char |

3. Collective Coommunication

3.1 sincronization

- Collective communication is a communication method that involves all processes in a communicator.
- In collective communication, synchronization among all process is required.
- All process cannot proceed until all processes reach the same point.
- To achieve this, MPI provides several collective communication functions.

3.2 MPI_Barrier

- **MPI_Barrier** is a collective communication function that synchronizes all processes in a communicator.
- All processes must call `MPI_Barrier` to ensure that all processes reach the same point before proceeding.
- It is often used to ensure that all processes have completed their previous tasks before moving on to the next step.
- The most basic usage of `MPI_Barrier` is to precise time measurement.
- If you do not call `MPI_Barrier` in all processes, the program will block and cannot proceed.
- `MPI_Barrier(MPI_Comm communicator);`

3.2 MPI_Barrier

```
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

printf("Rank %d: before barrier\n", rank);

sleep(rank);

MPI_Barrier(MPI_COMM_WORLD);
printf("Rank %d: after barrier\n", rank);
```

```
$ mpicc barrier.c -o barrier
$ mpirun -np 4 ./barrier
```

```
Rank 2: before barrier
Rank 3: before barrier
Rank 0: before barrier
Rank 1: before barrier
Rank 3: after barrier
Rank 1: after barrier
Rank 0: after barrier
Rank 2: after barrier
```


3.3 MPI_Bcast

- **MPI_Bcast** is a collective communication function that broadcasts data from one process to all other processes in a communicator.
- It is used to distribute data from a root process to all other processes.
- All processes in the communicator must call `MPI_Bcast` with the same parameters.

```
MPI_Bcast(  
    void* data,           // data buffer address  
    int count,           // number of elements in the buffer  
    MPI_Datatype datatype, // data type of the elements  
    int root,            // rank of the root process  
    MPI_Comm communicator // communicator  
);
```

- Root Process: Sends the data to all other processes.
- Other Processes: Receive the data from the root process.

3.3 MPI_Bcast

```
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

int data[10];
if (rank == 0) {
    for (int i = 0; i < 10; i++) {
        data[i] = i + 1;
    }
    printf("Rank 0: broadcasting data = [");
    for (int i = 0; i < 10; i++)
        printf(" %d", data[i]);
    printf(" ]\n");
}

MPI_Bcast(data, 10, MPI_INT, 0, MPI_COMM_WORLD);

printf("Rank %d: received data = [", rank);
for (int i = 0; i < 10; i++)
    printf(" %d", data[i]);
printf(" ]\n");
```

```
$ mpicc bcast.c -o bcast
$ mpirun -np 4 ./bcast
```

```
Rank 0: broadcasting data = [ 1 2 3 4 5 6 7 8 9
10 ]
Rank 0: received data = [ 1 2 3 4 5 6 7 8 9 10 ]
Rank 2: received data = [ 1 2 3 4 5 6 7 8 9 10 ]
Rank 3: received data = [ 1 2 3 4 5 6 7 8 9 10 ]
Rank 1: received data = [ 1 2 3 4 5 6 7 8 9 10 ]
```

3.3 MPI_Bcast

- We can implement MPI_Bcast wrapper using MPI_Send and MPI_Recv.

```
void my_bcast(void* data, int count, MPI_Datatype datatype, int root, MPI_Comm communicator) {
    int world_rank;
    MPI_Comm_rank(communicator, &world_rank);
    int world_size;
    MPI_Comm_size(communicator, &world_size);

    if (world_rank == root) {
        for (int i = 0; i < world_size; i++) {
            if (i != world_rank) {
                MPI_Send(data, count, datatype, i, 0, communicator);
            }
        }
    } else {
        MPI_Recv(data, count, datatype, root, 0, communicator, MPI_STATUS_IGNORE);
    }
}
```

- Q: Is this equivalent to MPI_Bcast?

3.3 MPI_Bcast

- A: **No**, it is less efficient than MPI_Bcast.
- This implementation has only one network communication link. The process with root rank sends data to all other processes one by one.
- MPI_Bcast uses Tree-based broadcast algorithm.
 1. The root process sends data to process 1.
 2. The root process sends data to process 2, process 1 sends data to process 3.
 3. The root process sends data to process 4, process 1 sends data to process 5, process 2 sends data to process 6, process 3 sends data to process 7.
- ...

3.3 MPI_Bcast

- Comparison of `MPI_Bcast` and `my_bcast`
- Average time of 10 trials

| Procs | Data Size | my_bcast (ms) | MPI_Bcast (ms) |
|-------|-----------|---------------|----------------|
| 16 | 40 | 0.008 | 0.009 |
| 16 | 400 | 0.022 | 0.009 |
| 16 | 4k | 0.026 | 0.010 |
| 16 | 40k | 0.096 | 0.018 |
| 16 | 400k | 0.355 | 0.084 |
| 16 | 4000k | 4.832 | 0.893 |
| 32 | 40 | 0.012 | 0.012 |
| 32 | 400 | 0.041 | 0.011 |
| 32 | 4k | 0.052 | 0.013 |
| 32 | 40k | 0.193 | 0.022 |
| 32 | 400k | 0.735 | 0.097 |
| 32 | 4000k | 7.447 | 0.937 |

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

3.4.1 MPI_Scatter

- **MPI_Scatter** is a collective communication function that distributes data from a root process to all other processes in a communicator.
- MPI_Bcast sends the same data to all processes, while MPI_Scatter sends different chunks of data to each process.

```
MPI_Scatter(  
    void* send_data,           // data buffer address to send  
    int send_count,           // number of elements to send to each process  
    MPI_Datatype send_datatype, // data type of the elements to send  
    void* recv_data,          // data buffer address to receive  
    int recv_count,           // number of elements to receive  
    MPI_Datatype recv_datatype, // data type of the elements to receive  
    int root,                 // rank of the root process  
    MPI_Comm communicator.    // communicator  
);
```

- send_count: the number of elements to send to each process.
- recv_count: the number of elements to receive from each process.

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

```
#define TOTAL_DATA 20
...
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

int send_data[TOTAL_DATA];
int recv_count = TOTAL_DATA / size;
int recv_data[recv_count];

if (rank == 0) {
    for (int i = 0; i < TOTAL_DATA; i++)
        send_data[i] = i;
    printf("Rank 0: Scattering data...\n");
}

MPI_Scatter(send_data, recv_count, MPI_INT,
           recv_data, recv_count, MPI_INT,
           0, MPI_COMM_WORLD);

printf("Rank %d received:", rank);
for (int i = 0; i < recv_count; i++)
    printf(" %d", recv_data[i]);
printf("\n");
```

```
// send_data == 5
```

```
$ mpicc scatter.c -o scatter
```

```
$ mpirun -np 4 ./scatter
```

```
Rank 0: Scattering data...
```

```
Rank 0 received: 0 1 2 3 4
```

```
Rank 1 received: 5 6 7 8 9
```

```
Rank 2 received: 10 11 12 13 14
```

```
Rank 3 received: 14 15 16 17 18
```

if `send_data` cannot divide by `size`, the last process will receive the remaining data.

```
// send_data == 6
```

```
Rank 0 received: 0 1 2 3 4 5
```

```
Rank 1 received: 6 7 8 9 10 11
```

```
Rank 2 received: 12 13 14 15 16 17
```

```
Rank 3 received: 18 19 20 -875497504 65535 20
```

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

3.4.2 MPI_Gather

- **MPI_Gather** is a collective communication function that collects data from all processes in a communicator and sends it to a root process.
- It is the reverse operation of `MPI_Scatter`.
- This is used in parallel sorting, parallel searching, and other parallel algorithms.

```
MPI_Gather(  
    void* send_data,           // data buffer address to send  
    int send_count,           // number of elements to send from each process  
    MPI_Datatype send_datatype, // data type of the elements to send  
    void* recv_data,          // data buffer address to receive  
    int recv_count,           // number of elements to receive from each process  
    MPI_Datatype recv_datatype, // data type of the elements to receive  
    int root,                 // rank of the root process  
    MPI_Comm communicator     // communicator  
);
```

- Except for the root process, pass `NULL` for `recv_data` is allowed.
- `recv_count` is the number of elements to receive from each process, not the total number of elements.

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

```
#define ITEMS_PER_PROC 2
...
int send_data[ITEMS_PER_PROC];
send_data[0] = rank * 2;
send_data[1] = rank * 2 + 1;

int recv_data[ITEMS_PER_PROC * size];

MPI_Gather(send_data, ITEMS_PER_PROC, MPI_INT,
           recv_data, ITEMS_PER_PROC, MPI_INT,
           0, MPI_COMM_WORLD);

if (rank == 0) {
    printf("Rank 0 gathered data: ");
    for (int i = 0; i < ITEMS_PER_PROC * size;
        i++)
        printf("%d ", recv_data[i]);
    printf("\n");
} else {
    printf("Rank %d sent data: %d %d\n", rank,
           send_data[0], send_data[1]);
}
```

```
$ mpicc gather.c -o gather
$ mpirun -np 4 ./gather
```

```
Rank 1 sent data: 2 3
Rank 3 sent data: 6 7
Rank 2 sent data: 4 5
Rank 0 gathered data: 0 1 2 3 4 5 6 7
```

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

3.4.3 MPI_Allgather

- MPI_Scatter and MPI_Gather conduct many-to-one or one-to-many communication.
- It is useful if you send data from multiple processes to multiple processes.
- **MPI_Allgather** is a collective communication function that collects data from all processes in a communicator and sends it to all other processes.
- It is like first MPI_Gather and then MPI_Bcast. Collect data by process rank order.

```
MPI_Allgather(  
    void* send_data,           // data buffer address to send  
    int send_count,           // number of elements to send from each process  
    MPI_Datatype send_datatype, // data type of the elements to send  
    void* recv_data,          // data buffer address to receive  
    int recv_count,           // number of elements to receive from each process  
    MPI_Datatype recv_datatype, // data type of the elements to receive  
    MPI_Comm communicator     // communicator  
);
```

- MPI_Allgather does not have a root process parameter.

3.4 MPI_Scatter, MPI_Gather, MPI_Allgather

```
#define ITEMS_PER_PROC 2
...
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

int send_data[ITEMS_PER_PROC];
send_data[0] = rank * 10;
send_data[1] = rank * 10 + 1;

int recv_data[ITEMS_PER_PROC * size];

MPI_Allgather(send_data, ITEMS_PER_PROC,
MPI_INT,
               recv_data, ITEMS_PER_PROC,
MPI_INT,
               MPI_COMM_WORLD);

printf("Rank %d received:", rank);
for (int i = 0; i < ITEMS_PER_PROC * size; i++)
    printf(" %d", recv_data[i]);
printf("\n");
```

```
$ mpicc allgather.c -o allgather
$ mpirun -np 4 ./allgather
```

```
Rank 2 received: 0 1 10 11 20 21 30 31
Rank 3 received: 0 1 10 11 20 21 30 31
Rank 0 received: 0 1 10 11 20 21 30 31
Rank 1 received: 0 1 10 11 20 21 30 31
```

3.5 MPI_Reduce, MPI_Allreduce

3.5.1 MPI_Reduce

- `reduce` is a basic concept in functional programming. It transforms a set of numbers into a smaller set of numbers.
 - `reduce([1, 2, 3, 4, 5], sum) = 15`
 - `reduce([1, 2, 3, 4, 5], multiply) = 120`
- Collect distributed data and apply a reduction operation is a tough task. However, MPI provides a simple interface to do this.
- **MPI_Reduce** is a collective communication function that collects data from all processes in a communicator to the root process, and applies a reduction operation to the data.

3.5 MPI_Reduce, MPI_Allreduce

```
MPI_Reduce(  
    void* send_data,          // data buffer address to send  
    void* recv_data,         // data buffer address to receive  
    int count,               // number of elements to send from each process  
    MPI_Datatype datatype,   // data type of the elements to send  
    MPI_Op op,               // reduction operation to apply  
    int root,                // rank of the root process  
    MPI_Comm communicator    // communicator  
);
```

- MPI Reduction Operations:
 - **MPI_MAX** - maximum value
 - **MPI_MIN** - minimum value
 - **MPI_SUM** - sum of values
 - **MPI_PROD** - product of values
 - **MPI_LAND** - logical AND of values
 - **MPI_LOR** - logical OR of values
 - **MPI_BAND** - bitwise AND of values
 - **MPI_BOR** - bitwise OR of values
 - **MPI_MAXLOC** - maximum value and its rank

3.5 MPI_Reduce, MPI_Allreduce

```
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

int value = rank + 1;
int sum;

MPI_Reduce(&value, &sum, 1, MPI_INT, MPI_SUM, 0,
MPI_COMM_WORLD);

double average = (double)sum / size;
if (rank == 0) {
    printf("Rank %d: sum = %d, avg = %.2f\n",
rank, sum, average);
} else {
    printf("Rank %d: sum = %d, avg = %.2f\n",
rank, sum, average);
}
```

```
$ mpicc reduce.c -o reduce
```

```
$ mpirun -np 4 ./reduce
```

```
Rank 1: sum = 4197236, avg = 1049309.00
```

```
Rank 0: sum = 10, avg = 2.50
```

```
Rank 2: sum = 4197236, avg = 1049309.00
```

```
Rank 3: sum = 4197236, avg = 1049309.00
```

3.5 MPI_Reduce, MPI_Allreduce

3.5.2 MPI_Allreduce

- **MPI_Allreduce** is a collective communication function that collects data from all processes in a communicator, applies a reduction operation to the data, and distributes the result to all processes.
- It is similar to `MPI_Reduce`, but the result is available to all processes, not just the root process.
- This is useful when all processes need to know the result of the reduction operation.

```
MPI_Allreduce(  
    void* send_data,          // data buffer address to send  
    void* recv_data,         // data buffer address to receive  
    int count,               // number of elements to send from each process  
    MPI_Datatype datatype,   // data type of the elements to send  
    MPI_Op op,               // reduction operation to apply  
    MPI_Comm communicator    // communicator  
);
```

3.5 MPI_Reduce, MPI_Allreduce

```
int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

int value = rank + 1;
int total_sum = 0;

MPI_Allreduce(&value, &total_sum, 1, MPI_INT,
MPI_SUM, MPI_COMM_WORLD);

double average = (double)total_sum / size;
printf("Rank %d: total sum = %d, average =
%.2f\n", rank, total_sum, average);
```

```
$ mpicc allreduce.c -o allreduce
$ mpirun -np 4 ./allreduce
```

```
Rank 2: total sum = 10, average = 2.50
Rank 1: total sum = 10, average = 2.50
Rank 0: total sum = 10, average = 2.50
Rank 3: total sum = 10, average = 2.50
```


4. References

4.1 MPI Reference

- [MPI「超」入門（C 言語編） - 東京大学情報基盤センター](#)
- [並列プログラミング入門](#)
- [MPI Tutorial](#)